

11-411/11-611 Homework Assignment 4

NLP Teaching Staff

Due: February 23, 2023

1 Introduction

In this homework, you will be building your first Language Model. You will be expected to build a N-gram Language model. You will also implement Laplace Smoothing (a Lazy version) to account for unknown words. The assignment is broken down into following subsections .

2 Learning Objectives

Language modeling is one of the most important concepts in natural language processing. As you learned in class, LMs assign a probability score to a sequence of words (or characters, or sub-words but we will limit ourselves to words!)

Today, language models have become ubiquitous: they are widely used in speech recognition, optical character recognition, machine translation and probably any other NLP task that you can imagine.

Thanks to language models, [insert your favorite voice assistant] can differentiate when you want an “ice cream” and when you just want to scream (“I scream”).

In this assignment, you will be programming your own language model powered by n-grams.

3 Task 1: Programming (60 points)

Refer to the notebook that is provided as part of the handout. You will be downloading the required data files, along with the **utils.py** and **main.py**. **Do not edit these files.** After you are done coding, paste the functions and the classes you implemented in the **lm.py** that is part of the handout and then upload the **utils.py** and **lm.py** to the HW4 Programming Submission, without zipping them.

4 Task 2: Written (40 points)

Answer the following questions based on the code you've written. You can use the latex file in the handout to answer the questions. Upload the PDF to the HW4 Written Submission

4.1 N-Gram counts(10 points)

Train the language model on the data/bbc/business.txt dataset for $n = 2$ and $n = 3$. Then do the same for data/bbc/sports.txt dataset

1. How many unique 2-grams are present in the business data-set?
2. How many unique 3-grams are present in the business data-set?
3. How many unique 2-grams are present in the sports data-set?
4. How many unique 3-grams are present in the sports data-set?

4.2 Song Attribution (8 points)

You are scrolling through the top hits playlist on Spotify when you notice a new unknown song at the top. It's recorded by an anonymous artist but the lyrics sound uncannily similar to some other songs you have heard. You have narrowed it down to three artists but are unable to choose one: *it could be any of them!*

You go along the rest of the day thinking who could it be. You reach Posner Hall to attend an NLP 11-411/611 lecture and David is teaching language models. Wait: language models. It suddenly hits you: language models can help in this task!

Train tri-gram ($n=3$, smoothing= 0.1) language models on collections of song lyrics from three popular artists ('data/lyrics/') and use the model to score a new unattributed song.

Note In reality, perplexity should only be used to compare language models when they have the same vocabularies but we will relax that condition for this question.

1. What are the perplexity scores of the test lyrics against each of the language models? (6 points)
 - (a) Taylor Swift:
 - (b) Green Day:
 - (c) Ed Sheeran:
2. Who is most likely to be the lyricist? (2 points)

4.3 Introduction to Decoding and Text Generation(8 points)

Run the code provided in the notebook and fill in the answers below

1. For each of these phrases 's1' to 's3', what are the top five word candidates after the sequence? Remember to not include, the special tokens we added during training to indicate end-of-sentence and start-of-sentence. (6 points)
 - (a) s1:
 - (b) s2:
 - (c) s3:
2. Report any one of the generated sentence here. Which generation mode do you think is better and why? (2 points)

4.4 Comparison to a GPT (14 points)

Run the code provided in the notebook and fill in the answers below.

1. What is the perplexity of your LM model? (2 points)
2. What is the perplexity of the GPT-2 model? (2 points)
3. How might you reason about the differences in perplexity between these two models? Think about the parameters, size of the vocabulary, perplexity training data etc. used to build your N-gram language model, compared to that of GPT-2, and how might this impact their respective performance? (6 points)
GPT-2 perplexity
4. What are some of the trade-offs between using a simpler model like your N-gram language model versus a more complex model like GPT-2, and how might these trade-offs affect their performance on different tasks? (4 points)

5 Deliverables

This assignment will be submitted via Gradescope in two parts. Upload both deliverables to gradescope.

1. **Coding.** Paste the functions and classes from the notebook into a file called `lm.py`, which is part of the handout. Upload `utils.py` and `lm.py` without zipping them.
2. **Written.** Submit answers to the questions as a PDF file.

Submissions are due February 23, 2023 at 11:59pm EST via Gradescope.